


Support Vector Machines

Warning!

- This is just a first introduction to SVMs – very basic version of SVM introduced
- We will discuss more on the advanced techniques later this semester
- Strong assumptions in this case


The Classification Problem



- Binary classification can be viewed as the task of separating classes in feature space
- w is the slope of the line
- b is the intercept

Linear Separators


- Which is the best linear separator?



- Depends on the goal
- Goal is to classify **accurately** and **generalize** to new examples.

Notion of Margins


Many different hyperplanes $\mathbf{w}'\mathbf{x} = b$ can classify the data. Which one will work best?



The hyperplane that maximizes the separation between the two classes (*the margin*)


Margins

- Distance from example \mathbf{x}_i to the separator is $r = \frac{\mathbf{w}^T \mathbf{x}_i + b}{\|\mathbf{w}\|}$
- Examples closest to the hyperplane are **support vectors**.
- **Margin** ρ of the separator is the distance between support vectors.



Intuition of a Margin

- Consider points A, B, and C
- We are quite confident in our prediction for A because it is far from the decision boundary.
- In contrast, we are not so confident in our prediction for C because a slight change in the decision boundary may flip the decision.



Given a training set, we would like to make all predictions correct and confident! This leads to the concept of margin.

Notation

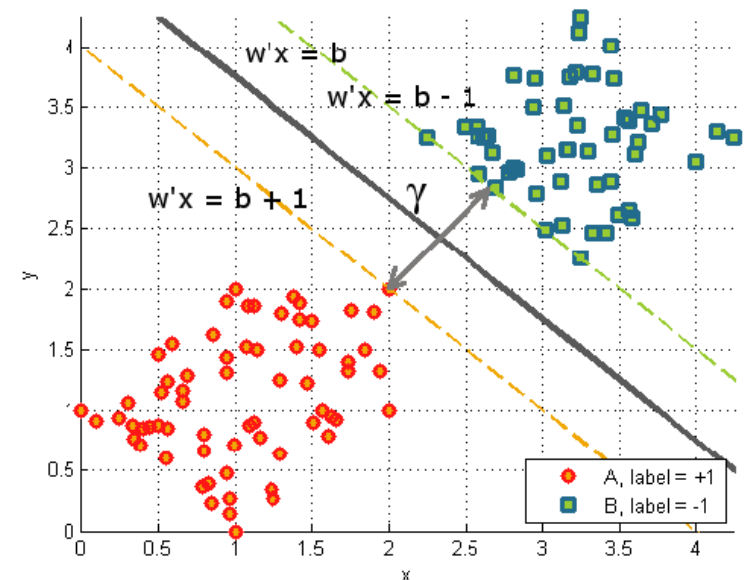
We denote the classifier, $f(\mathbf{x}) = \mathbf{w}'\mathbf{x} - b$, for all $\mathbf{x} \in \mathbb{R}^n$

assume supporting hyperplanes $\mathbf{w}'\mathbf{x} - b = \pm 1$

distance between the two hyperplanes is the *margin*, $\gamma = 2$



To achieve *scale invariance* divide the classifier by $\|\mathbf{w}\|_2$. Then the supporting hyperplanes are $\hat{\mathbf{w}}'\mathbf{x} - \hat{b} = \pm \frac{1}{\|\mathbf{w}\|_2}$.

margin is $\gamma = \frac{2}{\|\mathbf{w}\|_2}$.



Why Max Margin?

- Minimizes generalization error. Works well on **Future data**
- Minimizes Complexity. **Fewer support vectors**
- Minimizes the capacity of the classifier. **Eliminates overfitting**



Max margin Classifier

- Given a linearly separable training set $S=\{(\mathbf{x}^{(i)}, y^{(i)}) : i=1, \dots, N\}$, we would like to find a linear classifier with maximum margin.
- This can be represented as an optimization problem.

$$\max_{\mathbf{w}, b, \gamma}$$

$$\text{subject to : } y^{(i)} \frac{(\mathbf{w} \cdot \mathbf{x}^{(i)} + b)}{\|\mathbf{w}\|} \geq \gamma, \quad i = 1, \dots, N$$

Nasty optimization problem! Let's make it look nicer!

- Let $\gamma' = \gamma \cdot \|\mathbf{w}\|$, this is equivalent to

$$\max_{\mathbf{w}, b, \gamma'} \frac{\gamma'}{\|\mathbf{w}\|}$$

$$\text{subject to : } y^i (\mathbf{w} \cdot \mathbf{x}^i + b) \geq \gamma', \quad i = 1, \dots, N$$

Max margin Classifier

- Note that rescaling \mathbf{w} and b by $(1/\gamma')$ will not change the classifier, we can thus further reformulate the optimization problem

$$\begin{aligned} \max_{\mathbf{w}, b} \quad & \frac{\gamma'}{\|\mathbf{w}\|} \\ \text{subject to: } & y^i (\mathbf{w} \cdot \mathbf{x}^i + b) \geq \gamma', \quad i = 1, \dots, N \end{aligned}$$



$$\begin{aligned} \max_{\mathbf{w}, b} \quad & \frac{1}{\|\mathbf{w}\|} \quad (\text{or equivalently } \min_{\mathbf{w}, b} \|\mathbf{w}\|^2) \\ \text{subject to: } & y^i (\mathbf{w} \cdot \mathbf{x}^i + b) \geq 1, \quad i = 1, \dots, N \end{aligned}$$

Maximizing the geometric margin is equivalent to minimizing the magnitude of \mathbf{w} subject to maintaining a functional margin of at least 1

Solving the problem

$$\begin{array}{l} \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{subject to : } y^i (\mathbf{w} \cdot \mathbf{x}^i + b) \geq 1, \quad i = 1, \dots, N \end{array}$$

- This results in a ***quadratic optimization problem*** with *linear inequality constraints*.
- This is a well-known class of mathematical programming problems for which several (non-trivial) algorithms exist.
 - One could solve for \mathbf{w} using any of these methods
- We will see that it is useful to first formulate an equivalent dual optimization problem and solve it instead
 - This requires a bit of machinery

Constrained Optimization

- To solve the following optimization problem

$$\min_{\mathbf{x}} f(\mathbf{x}) \quad \text{subject to} \quad g_i(\mathbf{x}) \leq 0 \quad \text{for } i = 1, \dots, m$$

- Consider the following function known as the Lagrangian

$$\mathcal{L}(x, \alpha) = f(\mathbf{x}) + \sum_i \alpha_i g_i(\mathbf{x})$$

- Under certain conditions it can be shown that for a solution x' to the above problem we have

$$f(x') = \underbrace{\min_x \max_{\alpha} \mathcal{L}(x, \alpha)}_{\text{Primal form}} = \underbrace{\max_{\alpha} \min_x \mathcal{L}(x, \alpha)}_{\text{Dual form}}$$

$$\text{subject to } \alpha_i \geq 0$$

Dual Problem

$$\text{Minimize } \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{subject to: } 1 - y^i (\mathbf{w} \cdot \mathbf{x}^i + b) \leq 0, \quad i = 1, \dots, N$$

- The Lagrangian is

$$L(\mathbf{w}, b, \boldsymbol{\alpha}) = \frac{1}{2} \mathbf{w} \cdot \mathbf{w} + \sum_{i=1}^N \alpha_i \{1 - y^i (\mathbf{w} \cdot \mathbf{x}^i + b)\}, \text{ subject to } \alpha_i \geq 0$$

- We want to solve $\max_{\boldsymbol{\alpha}} \min_{w, b} \mathcal{L}(w, b, \boldsymbol{\alpha}) \quad s.t. \quad \alpha_i \geq 0$

- Setting the gradient of \mathcal{L} w.r.t. \mathbf{w} and b to zero, we have

$$\mathbf{w} - \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i = 0 \quad \Rightarrow \quad \mathbf{w} = \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i$$

$$\sum_{i=1}^N \alpha_i y^i = 0$$

Dual Problem

If we substitute $\mathbf{w} = \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i$ to \mathcal{L} , we have

$$\begin{aligned} L(\boldsymbol{\alpha}) &= \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \sum_{i=1}^N \alpha_i \{y^i (\mathbf{w} \cdot \mathbf{x}^i + b) - 1\} \\ &= \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle - \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle - b \sum_{i=1}^N \alpha_i y^i + \sum_{i=1}^N \alpha_i \\ &= \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle \end{aligned}$$

- Note that $\sum_{i=1}^N \alpha_i y^i = 0$
- This is a function of α_i only

Dual Problem

- The new objective function is in terms of α_i only
- It is known as the dual problem: if we know all α_i , we know \mathbf{w}
- The original problem is known as the primal problem
- The objective function of the dual problem needs to be maximized!
- The dual problem is therefore:


$$\max L(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j < \mathbf{x}^i \cdot \mathbf{x}^j >$$

subject to $\alpha_i \geq 0, i = 1, \dots, n,$



Properties of α_i when we introduce the Lagrange multipliers

$$\sum_{i=1}^N \alpha_i y^i = 0$$



The result when we differentiate the original Lagrangian w.r.t. b

- Note that there is only one constraint as against N in the original formulation
- Less number of variables

Dual Problem

$$\max L(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i \cdot \mathbf{x}^j \rangle$$

$$\text{subject to } \alpha_i \geq 0, i = 1, \dots, n, \quad \sum_{i=1}^N \alpha_i y^i = 0$$

- This is also quadratic programming (QP) problem
 - A global maximum of α_i can always be found

- \mathbf{w} can be recovered by $\mathbf{w} = \sum_{i=1}^N \alpha_i y^i \mathbf{x}^i$

- b can also be recovered as well (wait for a bit)

Characteristics of the Solution

- Many of the α_i are zero
 - \mathbf{w} is a linear combination of only a small number of data points
- In fact, optimization theory requires that the solution to satisfy the following KKT conditions:

$$\alpha_i \geq 0, i = 1, \dots, n,$$

$$y^i \left(\sum_{j=1}^N \alpha_j y^j < \mathbf{x}^j \cdot \mathbf{x}^i > + b \right) \geq 1$$

Functional margin ≥ 1

$$\alpha_i \{ y^i \left(\sum_{j=1}^N \alpha_j y^j < \mathbf{x}^j \cdot \mathbf{x}^i > + b \right) - 1 \} = 0$$

α_i is nonzero only when
functional margin = 1

- \mathbf{x}_i with non-zero α_i are called support vectors (SV)
 - The decision boundary is determined only by the SV
 - Let t_j ($j=1, \dots, s$) be the indices of the s support vectors. We can write

$$\mathbf{w} = \sum_{j=1}^s \alpha_{t_j} y^{t_j} \mathbf{x}^{t_j}$$

Solve for b

- Note that we know that for support vectors the functional margin = 1
- We can use this information to solve for b
- We can use any support vector to achieve this

$$y^i \left(\sum_{j=1}^s \alpha_{t_j} y^{t_j} < \mathbf{x}^{t_j} \cdot \mathbf{x}^i > + b \right) = 1$$

- A numerically more stable solution is to use all support vectors (details in the book)

Classifying new examples

- For classifying with a new input \mathbf{z}


- Compute $\mathbf{w}^T \mathbf{x} + b = \sum_{j=1}^s \alpha_{t_j} y^{t_j} \langle \mathbf{x}^{t_j} \cdot \mathbf{x} \rangle + b$ and classify \mathbf{z} as positive if the sum is positive, and negative otherwise

- Note: \mathbf{w} need not be formed explicitly, rather we can classify \mathbf{z} by taking a weighted sum of the inner products with the support vectors

(useful when we generalize from inner product to kernel functions later)

Support vectors

Only points, \mathbf{x}_i , that lie on the supporting hyperplanes have $\alpha_i > 0$. These are called the **support vectors**. Complexity of the solution only depends on the number of **support vectors**.




Recall that \mathbf{w} is a linear combination of training data

$$\mathbf{w} = \sum_{i=1}^N y_i \alpha_i \mathbf{x}_i = \sum_{\text{support vectors}} y_i \alpha_i \mathbf{x}_i$$

Support Vectors


Learned model will not change if we delete all the data




except the support vectors.

Geometric Perspective

Maximizing margin is equivalent to *maximizing distance between the two closet points on the convex hulls of the two sets.*



Geometric Perspective (2)



Summary

- We demonstrated that we prefer to have linear classifiers with large margin.
- We formulated the problem of finding the maximum margin linear classifier as a quadratic optimization problem
- This problem can be solved by solving its dual problem, and efficient QP algorithms are available.
- Problem solved?
- How about non-linear data? – Kernels
- How about noise? – Soft Margin SVMs